

SUMMARY:

AI Engineer specializing in Generative AI, LLM/RAG systems, and enterprise AI solutions. Experienced in building end-to-end AI pipelines using LangChain, LangGraph, FAISS, Pinecone, Ollama, and Python across computer vision (YOLO, OpenCV) and NLP domains, deployed on AWS and Docker. Currently engineering enterprise-grade applications on SAP BTP using SAP CAP (Node.js), CDS, SAP HANA Cloud, SAP Build Process Automation, and Master Data Management — combining deep AI expertise with full-stack enterprise development.

PROFESSIONAL EXPERIENCE:

AI Engineer Intern	Abhiyanta India Solutions	June 2026 – Present
- Architected and built an end-to-end Master Data Repository (MDR) and Governance Platform on SAP BTP using SAP CAP (Node.js), CDS, and SAP HANA Cloud, supporting 15+ enterprise master data records for centralized enterprise data governance and downstream S/4HANA syndication. - Designed a metadata-driven, table-agnostic architecture with runtime SQL/query builders, enabling configurable master data entities and validation rules across 20+ master data endpoints without backend code changes. - Implemented bulk Excel processing and master data validation via custom Node.js REST endpoints, including fuzzy duplicate detection using Levenshtein, Sorensen-Dice, and Soundex algorithms for in-batch and database-level duplicate resolution. - Integrated SAP Build Process Automation (BPA) for multi-stage approval workflows and configured XSUAA/AppRouter security, RBAC, OData V4, audit tracking, and draft/version management for governed master data processing. - Engineered CEO-Tower, an enterprise Generative AI and financial intelligence platform using Python, FastAPI/Flask, local LLMs (Llama.cpp), RAG pipelines, and Supabase; integrated Zoho Books APIs for real-time financial analytics, exception detection, and AI-driven proposal generation, containerized with Docker.		
Machine Learning Engineer	Automation Teknix	November 2025 – March 2026
- Constructed an end-to-end industrial object detection pipeline using YOLO on AWS SageMaker, covering data annotation, preprocessing, augmentation, model training, and evaluation across 1,000+ labeled images to enable automated defect detection and quality inspection. - Optimized YOLO models via TensorFlow Lite quantization and deployed within an offline Android application using OpenCV, enabling real-time on-device object detection and camera-based video processing without internet connectivity.		
Software Engineer (Intern)	Jnana Prabodhini	January 2025 – June 2025
Developed offline and online Virtual Kiosk applications using AngularJS, implementing JSON-based smart content routing, packaged .exe deployment, YouTube fallback integration, and a fully responsive cross-platform UI supporting multiple display environments.		

EDUCATION:

P.G Diploma in (AI & ML)	MITWPU	July 2024 – May 2025
Relevant Coursework: Machine Learning, Deep Learning, AI, NLP, Computer Vision, Neural Networks, Statistical Analysis		
BSc Computer Science	Abasaheb Garware College	August 2020 – June 2023

TECHNICAL SKILLS:

Programming:	Python, SQL, JavaScript, Node.js, C, R
AI/ML & GenAI:	Machine Learning, Deep Learning, Generative AI, LLMs, RAG, Agentic AI, NLP, Computer Vision, Prompt Engineering, TensorFlow, OpenCV, FastAPI
LLM & Retrieval:	LangChain, LangGraph, Ollama, FAISS, Pinecone, Vector Databases, Semantic Search
Cloud & Enterprise:	AWS (EC2, S3, Lambda, SageMaker, IAM, VPC), Docker, Docker Compose, GitHub Actions, CI/CD
SAP & Enterprise:	SAP BTP, SAP CAP, CAP Node.js, CDS, SAP HANA Cloud, SAP Build Process Automation, OData V4, REST APIs, XSUAA, AppRouter, Master Data Management

PROJECTS:

Multimodal RAG Chatbot – Ollama, LangGraph, FAISS
• Built a multimodal RAG pipeline using Ollama (Llama 2 7B), Nomic Embed Text embeddings, and FAISS vector search, enabling context-grounded semantic retrieval across PDF documents and OCR-extracted image data for accurate, hallucination-reduced responses. • Implemented stateful multi-turn conversational memory using LangGraph and SQLite, preserving session context across queries and grounding all responses in retrieved documents to minimize hallucinations.
Medical Chatbot (RAG, LangChain, GPT-4, Pinecone, AWS, Docker)
• Developed a production-ready RAG pipeline using Sentence Transformers and Pinecone vector database, enabling precise semantic retrieval across a 637-page medical knowledge base for accurate clinical query resolution. • Orchestrated LangChain with GPT-4 and structured prompt engineering to generate retrieval-augmented, clinically grounded responses, ensuring answer accuracy and traceability to source documents.

CERTIFICATION:

Oracle Cloud Infrastructure 2025 Certified Generative AI Professional | LLM, RAG